How would you push these data sets assume each new file is **incremental** from S3. Assume that there is a central data warehouse which the BI tool connects to (not S3).

As explain in the previous question's answer, we will be doing the following:

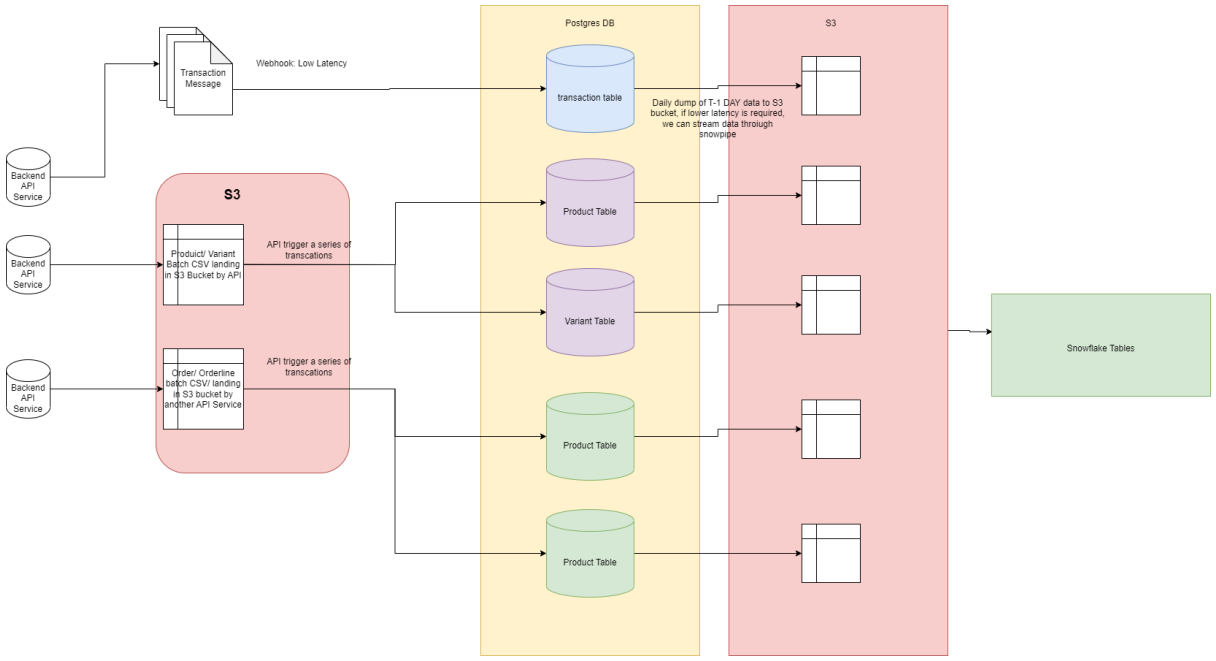

To briefly re-describe the solution that I have proposed in Q1, batch data from product/ variant is likely to come from a different business flow and API service, they will be ingested into Postgres through transactions. Upon landing the data in S3, we get an event trigger whenever a new batch csv arrives in the s3 bucket. This will trigger an Lambda Function to ingest the data into Postgres through transactions.

The same can be said about order/ orderline.

However, the Postgres layer will be strictly for backend usage and to be used by end users through an front end layer. Hence for subsquent BI analysis, we assume that BI is probably using cloud based datawarehouse tools such as Snowflake. We will do a periodic T-1 dump of data from Postgres layer back to S3. Then using the COPY INYO function that is triggered by another AWS lambda function, we will then

populate the ODS staging tables as described in Question 1. The final OLAP layers will be built in snowflake and BI tools can connect to snowflake tables to perform dashboarding.